

Image : filtering in the spatial domain

Image and signal processing

Samia Bouchafa-Bruneau

M1 E3A - UEVE/Upsay

- **Two kind of filters :**

- ① Linear Filters :

- Consists of a convolution between the input image and a given filter (convolution kernel) in the spatial domain. The spatial filter can be considered as the impulse response of the frequency domain filter.
 - Filter frequency cut-off correspond to the filter dimension.

- ② Non linear Filters :

- Non linear function of the neighborhood.

2D Discrete Convolution

- **2D Convolution**

$$f(x, y) * g(x, y) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(\alpha, \beta) g(x - \alpha, y - \beta) d\alpha d\beta$$

$$f_e(x, y) * g_e(x, y) = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} f_e(m, n) g_e(x - m, y - n)$$

- Important property of convolution :

$$\Im \{f(x, y) * g(x, y)\} = F(u, v) \cdot G(u, v)$$

$$\Im \{f(x, y) \cdot g(x, y)\} = F(u, v) * G(u, v)$$

2D Discrete Convolution

- Example in 1D :

$$A \begin{bmatrix} 1 & 2 & 3 & 4 & 5 \end{bmatrix} \quad * \quad$$

1	2	3	4	5
3	2	1		
3	4	3		
	10			

$$B \begin{bmatrix} 1 & 2 & 3 \end{bmatrix}$$

1	2	3	4	5
	3	2	1	
	6	6	4	
	10	16	22	

- **Managing image edges**

Many possibilities (depending on the application) :

- ① Zero padding
- ② Reflection
- ③ Default

2D Discrete Convolution

- Example in 2D :

10	5	20	20	20
10	5	20	20	20
10	5	20	20	20
10	5	20	20	20
10	5	20	20	20

$$\begin{bmatrix} -1 & 1 \\ 0 & 1 \end{bmatrix}$$

*

-10	15	-15	0	0
-10	15	-10	20	20
-10	15	-10	20	20
-10	15	-10	20	20
-10	15	-10	20	20

We have to consider the symmetric of

$$\begin{bmatrix} -1 & 1 \\ 0 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 1 & -1 \end{bmatrix}$$

Introduction

Finite impulse response Filter properties

• Properties

- Symmetry : → for low pass filters generally.
- Separability :

$$g(x, y) = [f(x, y) * h_x(x)] * h_y(x) = f(x, y) * [h_x(x) * h_y(x)] \text{ with}$$
$$h_{xy}(x, y) = h_x(x) * h_y(x)$$

Example : $\frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \end{bmatrix} * \frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$

- Cascade Filters :
$$g(x, y) = f(x, y) * h_1(x, y) * h_2(x, y) * h_3(x, y) * \dots = f(x, y) * h(x, y)$$
 with $h(x, y) = h_1(x, y) * h_2(x, y) * h_3(x, y) * \dots$
- Linear Combination of filters : $f(x, y) * h_1(x, y) + f(x, y) * h_2(x, y) = f(x, y) * [h_1(x, y) + h_2(x, y)] = f(x, y) * h(x, y)$ with $h(x, y) = h_1(x, y) + h_2(x, y)$.

Example :
$$\begin{bmatrix} 1 & 1 & 1 \\ -1 & 2 & -1 \\ -1 & -1 & -1 \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 \\ -1 & 2 & -1 \\ 0 & 0 & 0 \end{bmatrix} + \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

Low pass filtering

- We use low pass filtering for :
 - Smoothing
 - Reducing noise
 - Blurring the image

Low pass filtering

Average filtering

- Example of average filters of size 3×3 and 5×5 :

$$(1/9) \times \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$$(1/25) \times \begin{array}{|c|c|c|c|c|} \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline \end{array}$$

- These filters are generally :

- normalized
- centered and
- symmetric.

Low pass filtering

Average filtering

- Results with a 5×5 and 7×7 average filter :

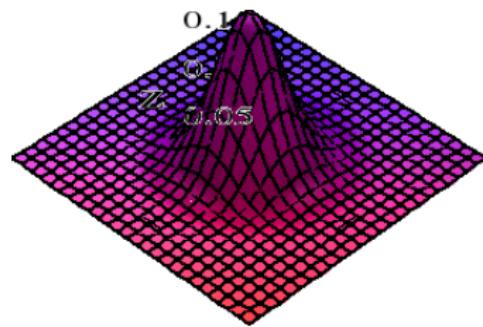


- Reduce noise but also details \rightsquigarrow blurring.

Low pass filtering

Gaussian Filter

$$G(x, y) = \left[\frac{1}{2\pi\sigma^2} \exp\left(-\frac{(x^2 + y^2)}{2\sigma^2}\right) \right]_{\sigma=1} = \frac{1}{2\pi} e^{-\frac{1}{2}x^2 - \frac{1}{2}y^2}$$

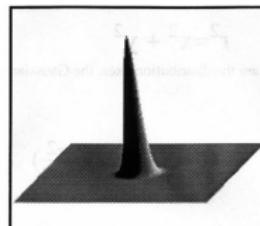


Low pass filtering

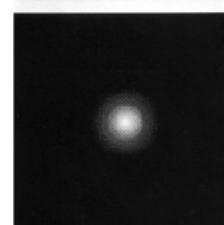
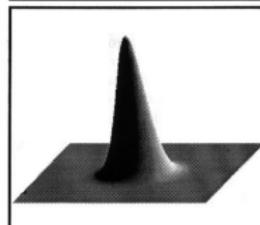
Gaussian Filter

- Varying the std :

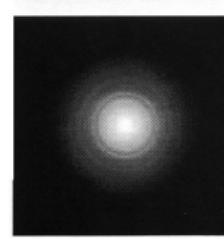
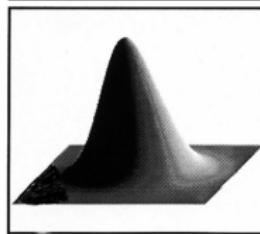
$$\sigma^2 = .25$$



$$\sigma^2 = 1.0$$



$$\sigma^2 = 4.0$$



Low pass filtering

Discrete Gaussian Filter

- Discrete Gaussian filter of

$$\text{size } n : W(i, j) =$$

$$\frac{1}{2\pi\sigma^2} \exp\left(-\frac{(i^2+j^2)}{2\sigma^2}\right) =$$

$$k \exp\left(-\frac{(i^2+j^2)}{2\sigma^2}\right) \text{ with}$$

$$k = \frac{1}{2\pi\sigma^2}. \text{ Then :}$$

$$\frac{W(i,j)}{k} = \exp\left(-\frac{(i^2+j^2)}{2\sigma^2}\right)$$

- Example for $n = 7$ and $\sigma^2 = 2$:

$$\sigma^2 = 2 \quad n = 7$$

		j						
		-3	-2	-1	0	1	2	3
i	-3	.011	.039	.082	.105	.082	.039	.011
	-2	.039	.135	.287	.368	.287	.135	.039
	-1	.082	.287	.606	.779	.606	.287	.082
	0	.105	.039	.779	1.000	.779	.368	.105
	1	.082	.287	.606	.779	.606	.287	.082
	2	.039	.135	.287	.368	.287	.135	.039
	3	.011	.039	.082	.105	.082	.039	.011

$\frac{W(1,2)}{k} = \exp\left(-\frac{1^2+2^2}{2*2}\right)$

Représentation entière, poser $k=91$ pour rendre ce coefficient unitaire

Low pass filtering

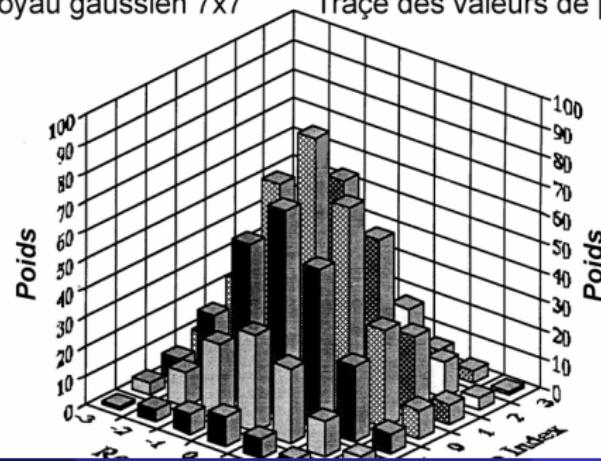
Discrete Gaussian Filter

1	4	7	10	7	4	1
4	12	26	33	26	12	4
7	26	55	71	55	26	7
10	33	71	91	71	33	10
7	26	55	71	55	26	7
4	12	26	33	26	12	4
1	4	7	10	7	4	1

$$\sum_{i=-3}^3 \sum_{j=-3}^3 W(i,j) = 1,115$$

Noyau gaussien 7x7

Traçé des valeurs de poids



Low pass filtering

Binomial Filters

- Discrete form of the Gaussian filter.
- Binomial filters and Gaussian filter are equivalent for high sizes of filters.
- Binomial Coefficients :
$$\binom{n}{k} = C_n^k = \frac{n!}{k!(n-k)!}$$
- Pascal's triangle :
$$\binom{n}{k} + \binom{n}{k+1} = \binom{n+1}{k+1}$$
- The filter coefficients are those of the Pascal's triangle :

n/k	0	1	2	3	4	5
1	1	1				
2	1	2	1			
3	1	3	3	1		
4	1	4	6	4	1	
5	1	5	10	10	5	1

Low pass filtering

2D Binomial Filters

- **1D binomial filter :**

1 $\begin{bmatrix} 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$

2 $\begin{bmatrix} 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 3 & 3 & 1 \end{bmatrix}$

etc.

- **For 2D binomial filters :** we use the separability property of filters

1 $\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$

2 $\begin{bmatrix} 1 & 3 & 3 & 1 \end{bmatrix} * \begin{bmatrix} 1 \\ 3 \\ 3 \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 3 & 3 & 1 \\ 3 & 9 & 9 & 3 \\ 3 & 9 & 9 & 3 \\ 1 & 3 & 3 & 1 \end{bmatrix}$

Low pass filtering

2D Binomial Filters

- Comparison between binomial filter and average filter of size 5×5 :



Low pass filters

Comparison between Average and Gaussian filters

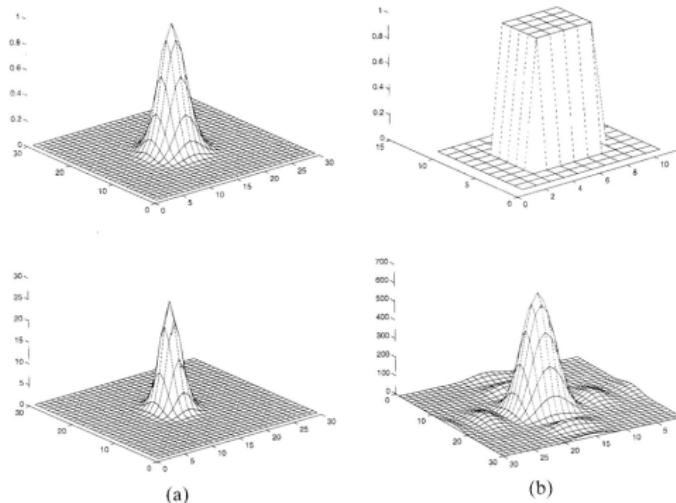


Figure 3.3 (a) The plot of a 5×5 Gaussian kernel of width 5 (top) and its Fourier transform (bottom). (b) The same for a mean-filter kernel.

- In low frequency, both filters have the same shape. In high frequency, Gaussian filter tend to 0 while the average filter rewind.

Low pass filtering

Nagao Filter

Image				
10	22	3	4	10
17	5	10	9	9
11	10	5	2	12
12	26	6	2	9
3	10	17	4	10

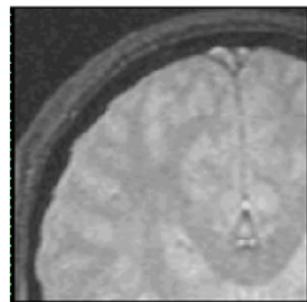
- Principle : replace each pixel value by the mean of the block with minimum std.

	Moyenne	Variance		Moyenne	Variance
	7.1	104.89		11.1	388.89
	9.2	227.32		11.3	360.02
	9.55	314.2		10.3	292.01
	11.2	553.56		7.7	109.37

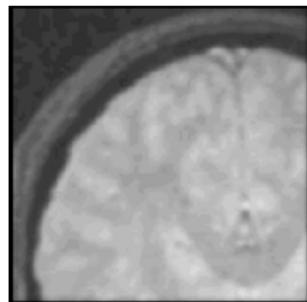
Nouvelle valeur du pixel = 7.1

Low pass filtering

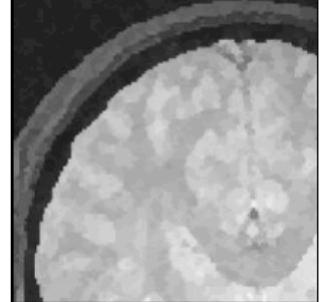
Nagao Filter - Example



Original image



After average filtering



After Nagao filtering

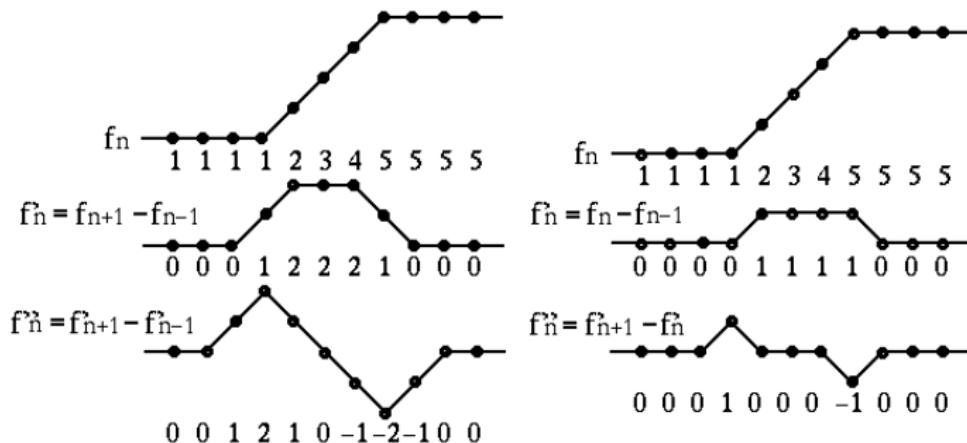
High pass filtering

- We use high pass filtering for :

- Transitions detection
- Edge detection
- Contrast increasing

- Two approaches :

- First order \rightarrow first derivative \rightarrow gradients
- Second order \rightarrow second derivative \rightarrow Laplacian



High pass filtering

Gradients

- **Definition :**

- Consider $f(x, y)$, then : $\nabla f = \begin{pmatrix} G_x \\ G_y \end{pmatrix} = \begin{pmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{pmatrix}$
- Gradient amplitude : $|\nabla f| = \sqrt{G_x^2 + G_y^2}$ or $|\nabla f| = |G_x| + |G_y|$
- Gradient orientation : $\theta = \arctan \left(\frac{G_y}{G_x} \right)$

High pass filtering

Gradients

- Gradient approximation :

$$\frac{\partial f}{\partial x} = \lim_{\Delta x \rightarrow 0} \frac{f(x+\Delta x, y) - f(x, y)}{\Delta x} \text{ or } \frac{\partial f}{\partial x} = \lim_{\Delta x \rightarrow 0} \frac{f(x, y) - f(x-\Delta x, y)}{\Delta x}$$
$$\frac{\partial f}{\partial y} = \lim_{\Delta y \rightarrow 0} \frac{f(x, y+\Delta y) - f(x, y)}{\Delta y} \text{ or } \frac{\partial f}{\partial y} = \lim_{\Delta y \rightarrow 0} \frac{f(x, y) - f(x, y-\Delta y)}{\Delta y}$$

- For $\Delta x = 1, \Delta y = 1 \Rightarrow$ this is equivalent to convolve the input image

with
$$\begin{bmatrix} 1 \\ -1 \end{bmatrix}$$
 and with
$$\begin{bmatrix} 1 & -1 \end{bmatrix}$$

- In fact, no reason to consider the right first derivative instead of the left first derivative or to consider that they are equal.

- We can chose the mean between them :

$$\frac{\partial f}{\partial x} = \frac{1}{2} \left[\frac{f(x+\Delta x, y) - f(x-\Delta x, y)}{\Delta x} \right] \text{ and } \frac{\partial f}{\partial y} = \frac{1}{2} \left[\frac{f(x, y+\Delta y) - f(x, y-\Delta y)}{\Delta y} \right]$$

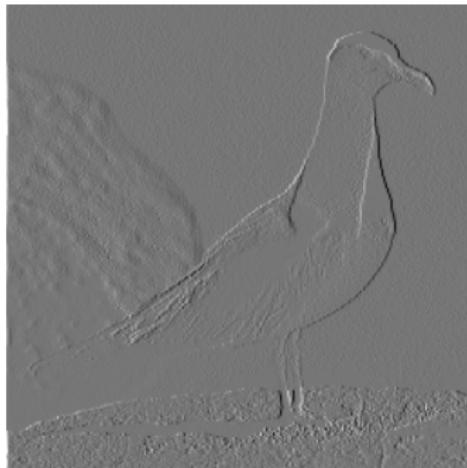
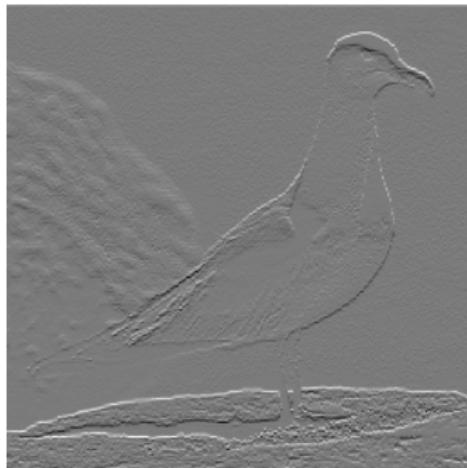
- For $\Delta x = 1, \Delta y = 1 \Rightarrow$ this is equivalent to convolve the input image

with
$$\begin{bmatrix} 1 & 0 & -1 \end{bmatrix}$$
 et avec
$$\begin{bmatrix} 1 \\ 0 \\ -1 \end{bmatrix}$$
.

High pass filtering

Gradients - Examples

- Example of x and y first derivative :



High pass filtering

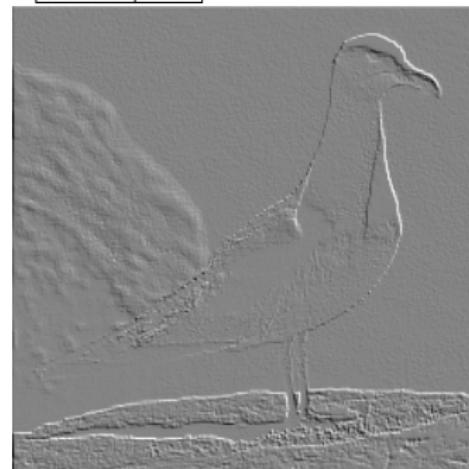
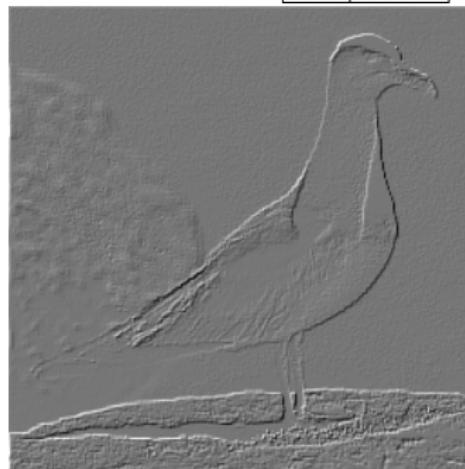
Gradients - Roberts kernels

- **Roberts kernels :**

$$\begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$

and

$$\begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$



High pass filtering

Gradients - Prewitt kernels

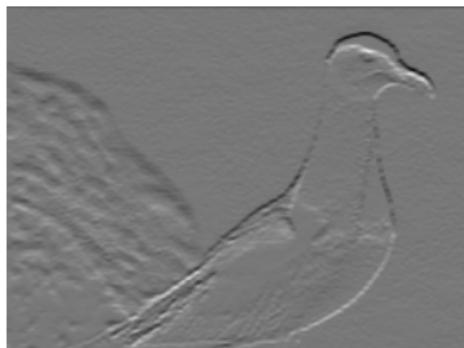
- **Prewitt kernels :**

- Average filter + first derivative

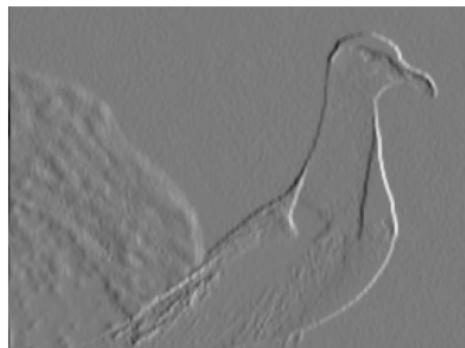
$$\begin{bmatrix} -1 \\ 0 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & 1 & 1 \end{bmatrix} = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

and

$$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$



(UEVE)



Spatial filtering

High pass filtering

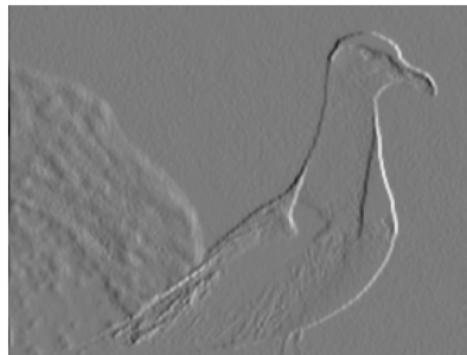
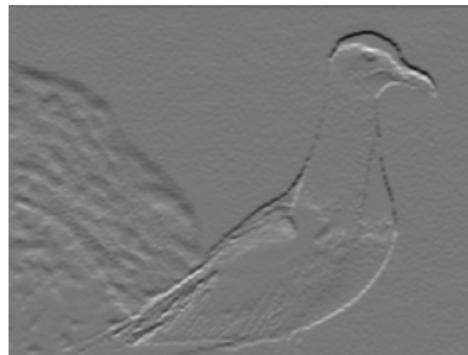
Gradients - Sobel kernels

- **Sobel Kernels :**

- Gaussian filter + first derivative

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} -1 & 0 & 1 \end{bmatrix} = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix} \text{ et}$$

$$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$



High pass filtering

Gradients - Kirsch kernels

- **Kirsch Kernels**

- Same principle than Prewitt but in all the possible edge's directions.
- For a 3×3 neighborhood \rightarrow 4 directions, 8 orientations.

-1	-1	-1
0	0	0
1	1	1
1	1	1
0	0	0
-1	-1	-1

-1	-1	0
-1	0	1
0	1	1
1	1	0
1	0	-1
0	-1	-1

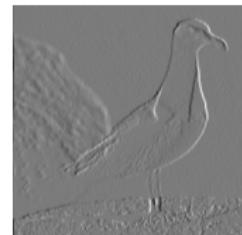
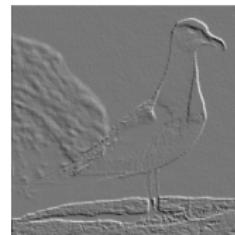
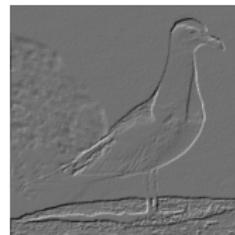
-1	0	1
-1	0	1
-1	0	1
1	0	-1
1	0	-1
1	0	-1

0	1	1
-1	0	1
-1	-1	0
0	-1	-1
1	0	-1
1	0	-1
1	1	0

High pass filtering

Gradients - Kirsch kernels

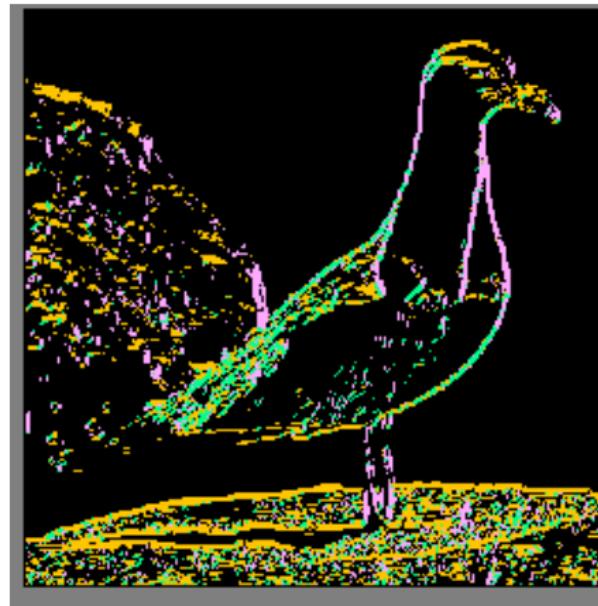
- Kirsch kernels : example



High pass filtering

Gradients - Kirsch kernels

- Kirsch kernels : example



High pass filtering

Laplacian

- Second derivative :

$$\frac{\partial^2 f}{\partial x^2} = \frac{\partial((f(x+1,y) - f(x,y))}{\partial x} = [f(x+2,y) - f(x+1,y)] - [f(x+1,y) - f(x,y)] = f(x+2,y) - 2f(x+1,y) + f(x,y)$$

- Convolve with the kernel :

1	-2	1
---	----	---

- Remark :

1	-2	1
---	----	---

 =

-1	1
----	---

 *

-1	1
----	---

- Laplacian operator :

$$\nabla^2 f = \left(\frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2} \right) \rightarrow$$

1	-2	1
---	----	---

 +

1		
-2		
1		

 =

0	1	0
1	-4	1
0	1	0

- Other forms :

0	-1	0
-1	4	-1
0	-1	0

or

-1	-1	-1
-1	8	-1
-1	-1	-1

 =

-1	0	-1
0	4	0
-1	0	-1

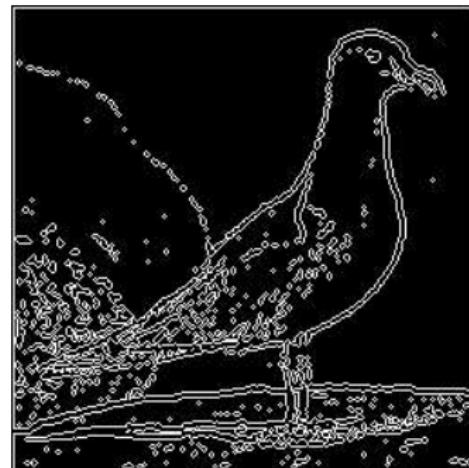
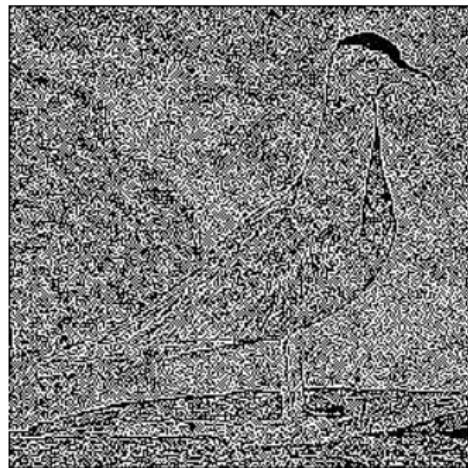
 +

0	-1	0
-1	4	-1
0	-1	0

High pass filtering

Laplacian - Example

- **Example :**



- **Advantages :** scalar \Rightarrow isotropic, just one convolution.

High pass filtering

Laplacian - other forms

- Average filter + second order derivative

$$\frac{1}{3} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & -2 & 1 \end{bmatrix} = \frac{1}{3} \begin{bmatrix} 1 & -2 & 1 \\ 1 & -2 & 1 \\ 1 & -2 & 1 \end{bmatrix}$$

- Gaussian filter + second order derivative

$$\frac{1}{4} \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} 1 & -2 & 1 \end{bmatrix} = \frac{1}{4} \begin{bmatrix} 1 & -2 & 1 \\ 2 & -4 & 2 \\ 1 & -2 & 1 \end{bmatrix}$$

High pass filtering

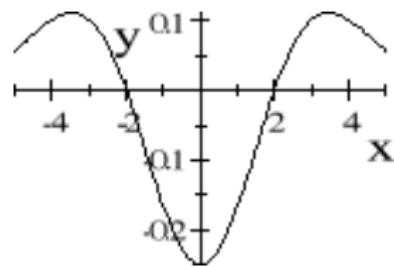
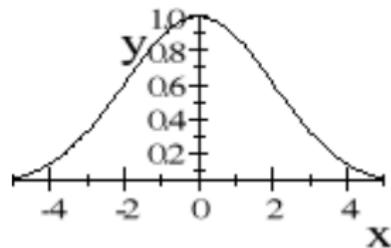
Laplacian - Marr-Hildreth Filter(LOG)

- Smoothing + Laplacian :

$$f(x, y) * G_\sigma * L = f * LoG \text{ (Laplacian of Gaussian)}$$

- In 1D : $G_\sigma(x) = \exp\left(-\frac{x^2}{2\sigma^2}\right) \rightarrow \frac{\partial^2 G}{\partial x^2} = \frac{-1}{\sigma^2} \left(1 - \frac{x^2}{\sigma^2}\right) \exp\left(-\frac{x^2}{2\sigma^2}\right)$

$$\left[G_\sigma(x) = \exp\left(-\frac{x^2}{2\sigma^2}\right) \right]_{\sigma=2} = 0 = \frac{1}{4} e^{-\frac{1}{8}x^2} \left(\frac{1}{4}x^2 - 1 \right)$$

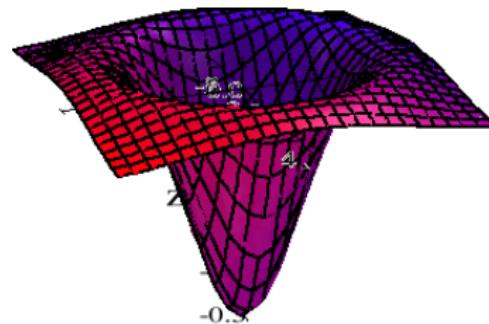


High pass filtering

Laplacian - Marr-Hildreth Filter(LOG)

- In 2D : $G_\sigma(x, y) = \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right) \rightarrow \text{LoG}(x, y) = \frac{\partial^2 G}{\partial x^2} + \frac{\partial^2 G}{\partial y^2} = \left(\frac{x^2+y^2-2\sigma^2}{\sigma^4}\right) \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right)$

$$\left[\left(\frac{x^2+y^2-2\sigma^2}{\sigma^4}\right) \exp\left(-\frac{x^2+y^2}{2\sigma^2}\right)\right]_{\sigma=2} = \frac{1}{16} e^{-\frac{1}{8}x^2-\frac{1}{8}y^2} (x^2 + y^2 - 8)$$



High pass filtering

Laplacian - Marr-Hildreth Filter(LOG)

- Example of LOG filter

0	0	-1	0	0
0	-1	-2	-1	0
-1	-2	16	-2	-1
0	-1	-2	-1	0
0	0	-1	0	0

pour $\sigma^2 = 4$

High pass filtering

Laplacian - Difference of Gaussians

- **Difference of Gaussians (DoG) :**

- The difference between two Gaussian filters is equivalent under some conditions to the LOG filter [Marr-Hildreth 80] :

$$DoG(x, y) = \frac{1}{2\pi\sigma_1^2} \exp\left(-\frac{x^2+y^2}{2\sigma_1^2}\right) - \frac{1}{2\pi\sigma_2^2} \exp\left(-\frac{x^2+y^2}{2\sigma_2^2}\right) \approx LoG(x, y) \text{ with } \sigma_1 > \sigma_2$$

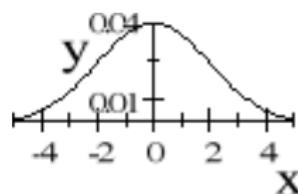
The two functions have the same zero-crossings if :

$$\sigma^2 = \frac{\sigma_1^2\sigma_2^2}{\sigma_1^2 - \sigma_2^2} \log\left(\frac{\sigma_1^2}{\sigma_2^2}\right)$$

High pass filtering

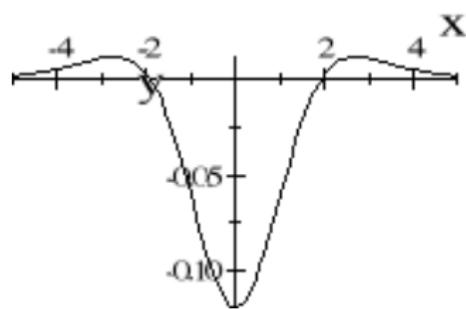
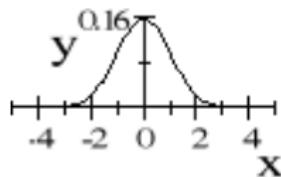
Laplacian - Difference of Gaussians

$$\left[G_\sigma(x) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2}{2\sigma^2}\right) \right]_{\sigma=2}$$



$$DoG(x) = \frac{1}{8\pi} e^{-\frac{1}{8}x^2} - \frac{1}{2\pi} e^{-\frac{1}{2}x^2}$$

$$\left[G_\sigma(x) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2}{2\sigma^2}\right) \right]_{\sigma=1}$$



- **Non-linear filters :**

- Non-linear function of the neighbors
- Require defining the non-linear operator + neighborhood size.

- **Example of non-linear filters :**

- Min → Each pixel value is replaced by the minimum value in the neighborhood.
 - Max → Each pixel value is replaced by the maximum value in the neighborhood.
 - Majority → Each pixel value is replaced by the most frequent value in the neighborhood.
 - Extreme → Each pixel value is replaced by the nearest between min and max value in the neighborhood.
 - etc.
- * Particular case : binary image → Min=AND Max=OR operators

Median Filter

- Median filter : each pixel value is replaced by the median value in the neighborhood → remove impulse noise

10	12	40	16	19	10
14	22	52	10	55	41
10	14	51	21	14	10
32	22	9	9	19	14
41	18	9	22	27	11
10	7	8	8	4	5

Liste non-ordonnée des valeurs de pixel

51
21
14
9
9
19
9
22
27

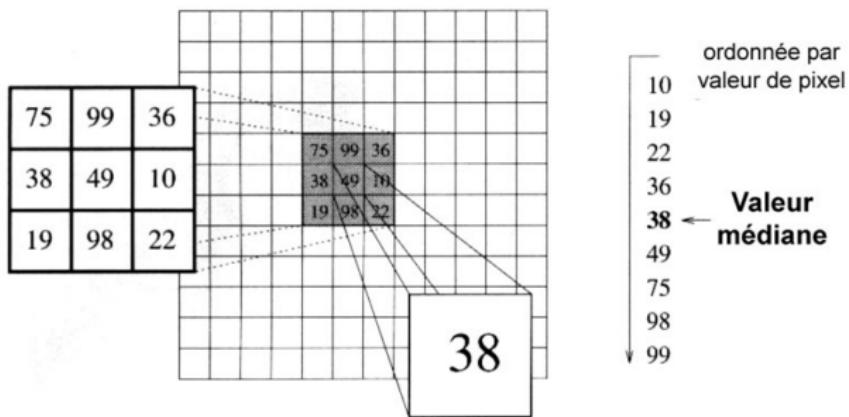
Liste ordonnée des valeurs de pixel

9
9
9
14
19
21
22
27
51

nouvelle valeur du pixel

Même nombre de valeurs au dessus et en dessous de la médiane

Median Filter



Median Filter

- Many possibilities concerning the selection of pixels in the neighborhood :

1	1	1
1	1	1
1	1	1

0	1	0
1	1	1
0	1	0

0	0	1	0	0
0	0	1	0	0
1	1	1	1	1
0	0	1	0	0
0	0	1	0	0

- Efficient for suppressing extreme values.
- If no optimization, the filter is time consuming.
- Details are suppressed .
- No new pixel values.

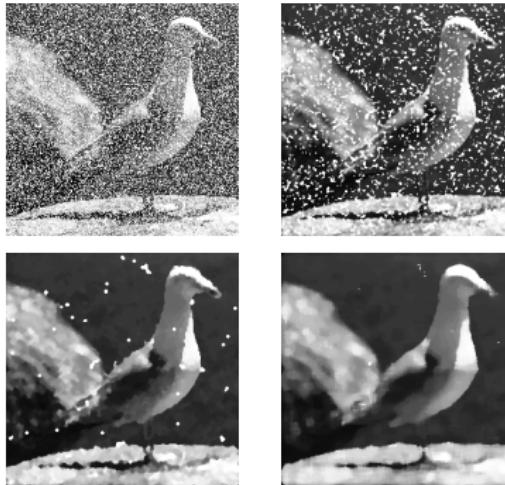
Median filter - Example 1

- Salt and Pepper noise of density 10%
- Noise could be suppressed if its density is less than half the filter size.



Median filter - Example 2

- Salt and Pepper noise of density 40%, 3×3 median filter, 3×3 median filter repeated 10 times and 9×9 median filter.



From contrast points to edges



Threshold =10



Threshold =100

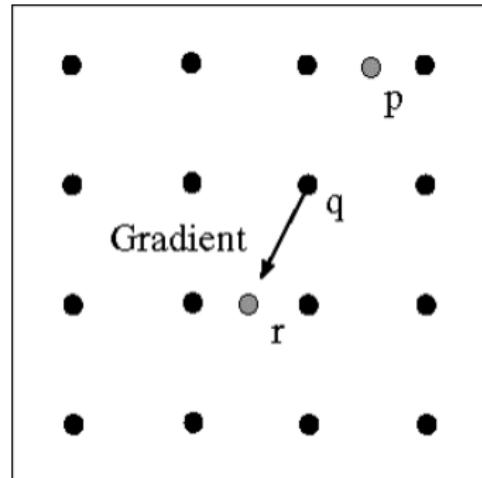
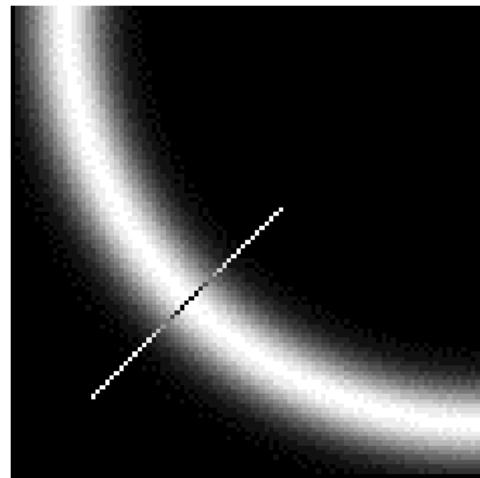


Threshold =200

The choice of a suitable threshold is delicate. Threshold too low : thick contours + noise. Threshold too high : discontinuities.

The Canny Edge detector

Non-maximum suppression

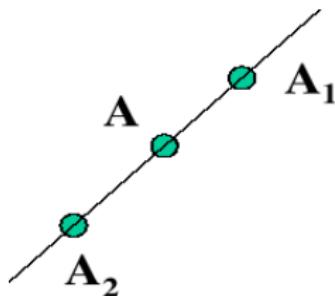


Check if pixel is local maximum along gradient direction (perpendicular to contour), select single max across width of the edge
→ requires checking interpolated pixels p and r

The Canny Edge detector

Eliminate points that are not local *extrema* in gradient direction

- Starting from each edge point A , consider its direct neighbors A_1 and A_2 in gradient direction.
- A is a local maximum if : $G(A) > G(A_1)$ and $G(A) > G(A_2)$



The Canny Edge detector

Hysteresis thresholding

Use a high threshold to start edge curves, and a low threshold to continue them.

- High threshold \rightarrow most reliable edges. Let f_1 be the binary resulting image.
- Low threshold \rightarrow Edges + noise. Let f_2 be the binary resulting image.
- Keep only points in f_2 that are connected to a non null point in f_1 .

The Canny Edge detector

Hysteresis thresholding



Original image



High threshold (strong edges)



Low threshold (weak edges)



Hysteresis threshold

The Canny Edge detector

- Filter image with derivative of Gaussian
- Find magnitude and orientation of gradient
- Non-maximum suppression :
 - Thin wide “ridges” down to single pixel width
- Linking and thresholding (hysteresis) :
 - Define two thresholds : low and high
 - Use the high threshold to start edge curves and the low threshold to continue them

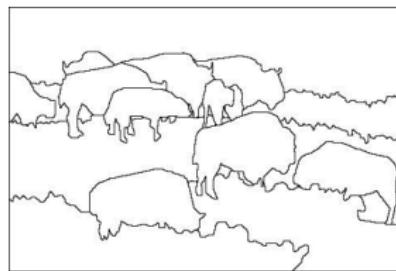
J. Canny, A Computational Approach To Edge Detection, IEEE Trans. Pattern Analysis and Machine Intelligence, 8 :679-714, 1986.

The Canny Edge detector

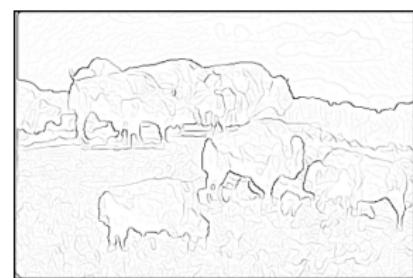
Edge detection is just the beginning...



Image



Human
segmentation



Gradient magnitude